

Converting a PDF textbook to be accessible

Timothy Prescott

Abstract

A guideline in the United States requires that PDFs created after April 2026 must be accessible to visually impaired users. We discuss how we updated a custom three semester calculus book to be accessible, including what we did, the challenges we faced, how we verified accessibility, and what remains to be done. We also discuss necessary adaptations to satisfy a common accessibility checker available in several LMSs.

1 Introduction & Setup

Since 2016, the University of North Dakota has used a custom 1150 page textbook for its three semester calculus sequence. The source for the textbook is at github.com/teepeeemm/APEXCalculusLT_Source, forked from github.com/APEXCalculus/APEXCalculus_Source while the compiled document is at arts-sciences.und.edu/academics/math/calc-1-texts.html.

Using a custom textbook has enabled students to download the PDF for free in addition to having the option to purchase the book for a fraction of the cost of other textbooks (currently \$25–\$30 per semester, when the author last checked).

The format of the book is to `\include` one `\part` for each semester, so that `\includeonly` allows us to have a one semester PDFs that we send to the university printer to create physical copies (the online PDF includes all three semesters). Within each included file, we have a few paragraphs introducing each chapter and then `\input` each section. Each section then ends with a custom command that inputs the exercises for that section.

Our university has been emphasizing the national requirement that all documents must be accessible according to the WCAG 2.1 AA guideline as of April 2026 (now delayed to 2027), see *TUGboat* 45:2 [?] for a historical overview of how we got to this point. Beginning in early 2025, we began exploring how to ensure that the textbook meets this guideline.

2 The Basics

We first needed to convert the document from using $\text{\LaTeX}$ to $\text{\Lua\LaTeX}$, which also necessitated adjusting a few font commands in the preamble, but not much else.

The $\text{\LaTeX}$ team has been reporting their progress developing tagging since at least *TUGboat* 40:2 [?] (see also tug.org/TUGboat/Contents/listkeyword.html#CatTAGAccessibility), building off of exist-

ing efforts. Thanks to their hard work, the vast majority of accessibility is enabled by using

```
\DocumentMetadata{
  lang=en,
  tagging=on,
  pdfstandard=ua-2,
% check-tagging-status,
  tagging-setup={
%   math/alt/use,
    math/setup=mathml-SE
  }
}
```

as the first command of the document (even before the `\documentclass`). We'll discuss the commented options as we proceed; other options are available by looking through `texdoc documentmetadata-support-code`. In addition to enabling tagging, the `\DocumentMetadata` command also allows us to adjust the PDF metadata of the compiled document. We also need to use the package `unicode-math`.

An all too common occurrence during compilation was a tagging related error, often that “The number of automatic begin and end hooks differ!” Our setup of using `\include` and `\input` allowed us to comment out various sections, isolating a single section that caused the error. We then examined the log file, found a “Package tagpdf Warning” regarding some aspect of our code, and created a minimal working example causing this warning and error.

With the example in hand, we began searching the internet for a fix. For packages, a good starting point was to view latex3.github.io/tagging-project/, `tagging-status/`, which reports the tagging compatibility of each class and package. Often, a linked issue would contain code snippets we could copy and paste into the preamble that would resolve the issue. We can also create a tagging status report in the log file by uncommenting the `check-tagging-status` line in our `\DocumentMetadata`.

The `tagging-status` website and log information have five categories:

compatible Everything should work for the package, and the $\text{\LaTeX}$ team wants to know if it doesn't. This is the most common entry for the website.

partially-compatible The package should work if you avoid issues described in the website.

currently-incompatible Double check the issue(s) in the website to see if you can make it work.

no-support / unsupported You need to find a new package. (As an aside, if a required package in this category is causing errors because you use `\DocumentMetadata` solely to adjust the

PDF metadata, you can use the `pdfmanagement` package instead of `\DocumentMetadata`.)

unchecked / unknown / unclassified The package has not been tested. Your odds are good if some other package loaded it; your odds are bad if it's a local package provided by a "friend".

Resolving errors not directly related to packages was more difficult. We often turned to `tex.stackexchange.com/` for help. Luckily for the reader (and the author), some errors were caused by a bug in the tagging code that was quickly fixed. (A quickly fixed bug was also occasionally true of errors relating to packages.)

3 Alt Text

With the bulk of the work handled automatically by `LuaLATEX`, the greatest part of the work requiring manual intervention was providing alternative text for the thousand images throughout the book. Alt text provides a verbal description of an image that a visually impaired person has difficulty seeing. Ideally, alt text should be

- brief (at most 120 characters or 2 sentences)
- objective
- not the sole source of necessary information (most users will not see the alt text)
- not begin with the word "image", so that a screen reader does not say "Image: image ..."
- not blindly copied from an AI output

`LATEX`'s tagging code allows us to specify alt text as an optional argument to the image:

- `\includegraphics[alt={...}]`
- `\tikz[alt={...}]`
- `\begin{tikzpicture}[alt={...}]`
- `\begin{picture}[alt={...}]`

We also have two other options to `alt={...}`: if the image is purely decorative and unnecessary for the reader, we can use the option `artifact` so that screen readers ignore it entirely. Less commonly, we can use `actualtext={...}` so that screen readers act as if there is no image, and instead read the text. This works best if the image is actually stylistically decorated text.

In our case, we were able to hire a student worker over the summer to supply the bulk of the alt text.

There are two contrasting problems arising from alt text. The first was succinctly stated by a colleague: "What alt text would you give for the snake lemma?" (see Figure 1). It is not reasonable to describe a commutative diagram with fourteen cells in at most 120 characters. On the other hand, there is no *mathematical* problem with having a longer




Figure 1: The conclusion of the snake lemma. Well, how would you write alt text for this? Image code adapted from `tex.stackexchange.com/a/440352/`.

Exercise: Describe this function, identifying limiting values, intercepts and their degree, critical and inflection points, and regions where the function is increasing, decreasing, and concave up or down.




Figure 2: How do you write alt text for this exercise without giving away the answer?

description. Accessibility guidelines suggest creating a short alt text pointing to a longer description elsewhere.

A bigger problem arises in early chapters in the calculus sequence when we ask students to verbally describe a graph, as in Figure 2. We chose to describe the graph using various synonyms ("the graph is a line from $(-2, 0.5)$ to a hollow dot at $(-1, 1)$, ..."), but acknowledge this is not an entirely satisfactory approach.

3.1 Other Issues Related to Alt Text

Many PDF checkers do not recognize PDF version 2.0 and instead check the PDF against version 1.7. A critical distinction for us is that PDF version 1.7 requires that formulas include alt text. We can accomplish this by uncommenting the `math/alt/use` key in the `\DocumentMetadata` command. This usually causes PDF screen readers to read the alt text and omit the more informative MathML tags, so we discourage its use unless a document is being judged on its adherence to the 1.7 standard.

`rmarkdown: \setkeys{Gin}{alt={...}}`

4 Tables

Tables need special consideration so that accessibility tools are able to identify the headings that describe

the data in the table. But first, it should be emphasized that if the sole purpose of the table is to arrange elements in a grid, then it's not actually a table and probably shouldn't be tagged as such (there are ways to tell the tagging code "this doesn't need to be tagged as a table", but life will probably be easier if it's not a table to begin with).

To identify the headings in a particular tabular, we use

```
\tagpdfsetup{
  table/header-rows={...},
  table/header-columns={...}
}
```

The default is an empty list specifying no header along that axis, so that the most common occurrence of a single header row only needs

```
\tagpdfsetup{table-header-rows={1}}
```

and can omit the `table-header-columns`. These declarations obey TeX's grouping rules, so that a header specification will last until the end of the current environment.

While `\multicolumn` is tagging ready, we cannot say the same for a `\multirow` cell, which needs to have

```
\tagpdfsetup{table/multirow={\langle number of rows \rangle}}
```

at the beginning of the cell.

We should also note that the table interface will likely change over the next few years. Follow `texdoc latex-lab-table` for the details.

5 Other Package Changes

We had originally formatted our headings using the `titlesec` package. This requires adjustments to properly tag, and was becoming too complicated for our use. We opted to instead use `\@startsec` (see `texdoc latex2e`, section 6.8). Another possibility is the new template mechanism (see `texdoc latex-lab-sec-template`).

We were originally using `enumitem` to adjust our `{enumerate}` environments, which is also not amenable to tagging. We are now using `enumext`, although there is also a templating alternative (see `texdoc blocks-doc`).

We also made heavy use of the `marginnote` package which, along with `marginfit`, is currently incompatible with tagging. Instead, we are now using `marginalia`.

6 Checking the Results

Having resolved all possible errors, we needed to verify that our result was indeed accessible. The obvious first step is to ensure a screen reader properly handles the document. Unfortunately, when it comes

Table 1: Successful validations examined

Profile	Rules	Checks
PDF/UA-2	1723	1314 0427
WTPDF 1.0 Accessibility	1719	1314 0423
WTPDF 1.0 Reuse	1706	1190 4349

to PDF version 2.0, the only currently viable PDF readers are Adobe, Foxit, and Firefox, while the only currently viable screen readers are to go with them are NVDA and JAWS on Windows and Orca on Linux. See tex.stackexchange.com/a/755945/ for the latest.

Another approach is to verify that our PDF is indeed version 2.0, and satisfies the UA-2 PDF standard. Unfortunately (again), there are not many PDF checkers that recognize version 2.0, and many will apply version 1.7 standards to a version 2.0 document, resulting in many claimed "issues" with a valid document. pdfa.org/accessible-math-in-pdf-finally/ lists several (but currently only five) checkers that support PDF version 2.0. We have experience with veraPDF (verapdf.org/) and PDFfix (pdfix.net/).

The L^AT_EX team uses veraPDF to validate their documents, so that using the same validation gives similar results (this is usually good, but can be bad if they have overlooked an issue). The main downside we found with veraPDF is that while it provides a GUI to initiate a validation and see the results, it reports the location of failed validation checks as paths within the PDF document, such as

```
root/
  document[0]/
    pages[65](167 0 obj PDPPage)/
      annots[0](11747 0 obj PDLINKAnnot)
```

VeraPDF also has a command line version of the program, allowing for faster (or even automated) checking of generated PDFs. After a minute examining the Calculus textbook, veraPDF presents an XML report that the text book complies with various profiles and a number of rules and checks for each profile, summarized in Table 1.

PDFfix uses a similar checker as veraPDF, but it also provides a graphical view of the PDF in question, which better locates errors. PDFfix also allows navigating the tree structure of the PDF, making it possible to locate errors given by veraPDF.

Our institution uses Ally to check PDFs posted to our LMS. Ally is one of the aforementioned PDF checkers that looks for version 1.7 and requires formulas to have alt text. For this reason, we try to avoid posting the textbook to the LMS.

Ally also requires that documents over two pages use headings. We normally accomplish this with

Table 2: Factor by which various statistics increase for compiling the tagged textbook.

Statistic	Factor
Time	6.7
RAM	4.4
Output size	1.3
strings	11
hash	11

`\section` commands, but if a document does not otherwise have them, it is possible to have the title also provide a heading, see `tex.stackexchange.com/a/758805/`.

For shorter documents, we can also examine the tagging structure of a PDF by using the command line tool `show-pdf-tags` or the website `ngpdf.com/`.

* Ally, UDOIT, SiteImprove, continual engine prep? * Adobe?

Accessible PDFs require more computing power to create. In Table 2, we report the factor by which various computational statistics increase for the textbook. The final two entries relate how much memory the compilation process allocates. This originally resulted in fatal compilation errors that $\text{\TeX}$ exceeded its capacity for the `number of strings` or the `hash size`, requiring us to modify the compilation command (see `tex.stackexchange.com/a/741777/`).

This is an example article for *TUGboat*, linked from `tug.org/TUGboat/location.html`.

7 Basic packages and hyperlinks

The standard `graphicx` and `xcolor` packages work fine for image and color handling, respectively, as does the `hyperref` package for active urls. *TUGboat* is produced using PDF files exclusively, but conversion from DVI is fine and supported.

Please use the standard `\url` command for urls:

```
\url{https://some/url} % yes
```

For *TUGboat*, we prefer not to print the `'https://'`, but to include the protocol in other cases, so we usually do `\def\url{\tburl}`, which takes care of that, while keeping the source using the standard `\url` command.

In any case, please do *not* hide hyperlinks via `\href`, as in:

```
\href{https://some/url}{some text} % no!
```

When printed on paper, the url will not be visible and readers will not know it exists.

8 Abbreviation macros and typing

The `ltugboat` class provides many abbreviation commands; here are a few of the most common:

```
\AllTeX \AMS \BibTeX \Cplusplus \CTAN \DVI
\HTML \LaTeXe \macOS \MathML \MF \PDF \PS
\TUB \TUG \tug \WEB \XeLaTeX \XeTeX \XML
```

A few other typing conventions:

- For an em-dash with our spacing (preferred to `---`): `\Dash`, with output `—` like this.
- For initialisms in all caps: `\acro{FRED}`, with output: `FRED`.
- A literal control sequence: `\cs{fred}`, with output: `\fred`.
- A syntactic metavariable: `\meta{fred}`, with output: `⟨fred⟩`.
- A title: `\titleref{Book of Fred}`, with output: *Book of Fred*.

We recommend using `\begin{verbatim}[] ... \end{verbatim}` for code blocks, since we prefer not to colorize code when printed. But if you want to have some font changes, our recommended settings for the `listings` package are in the `ltugguid` manual mentioned below.

Please put punctuation outside quotes, “like this”, unless the punctuation is actually part of the quoted material.

Also, please put punctuation after footnotes.

9 Figures

For *TUGboat*, the standard `figure` environment produces a column-width figure; this is desirable when at all possible. The `figure*` environment produces a full-width (across both columns) figure when needed. Analogously for `table` and `table*`.

Please put captions below figures, but above tables.

Don’t worry overmuch about figure placements, as they will likely change with editing.

10 References

For references to other issues of *TUGboat*, please use the format *volno:issno*, e.g., “*TUGboat* 32:1”.

The primary *TUGboat* style documentation is the `ltugguid` manual, available online at `ctan.org/pkg/tugboat` or locally with `texdoc tugboat`. For general CTAN package references, we recommend that form, using `/pkg/`. If you need to refer to a specific file on CTAN, use:

```
https://mirror.ctan.org/path
```

We recommend using $\text{\BIBTeX}$ (but don’t require it), with the *tugboat* $\text{\BIBTeX}$ style. The `biblio` directory on CTAN, `ctan.org/pkg/biblio`, provides files `tugboat.bib` with a complete bibliography of *TUGboat*, `texbook3.bib` with many common $\text{\TeX}$ -related books and articles, and plenty more.

Email `tugboat@tug.org` with any questions.

◇ Timothy Prescott
University of North Dakota
101 Cornell Street Stop 8376
Grand Forks ND, 58203
USA
`timothy dot prescott (at) UND`
`dot edu`
ORCID 0009-0009-0555-8074